

final-notebook

June 11, 2024

1 Predicting Recipe Rating based on Recipe Length

Name(s): Jahnavi Naik

Website Link: <https://jahnavi-naik.github.io/recipe-analysis/>

```
[ ]: import pandas as pd
import numpy as np
from pathlib import Path

import plotly.express as px
pd.options.plotting.backend = 'plotly'
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score

# from dsc80_utils import * # Feel free to uncomment and use this.
```

1.1 Step 1: Introduction

```
[ ]: #merging the datasets
recipes_fp = 'food_data/RAW_recipes.csv'
review_fp = 'food_data/RAW_interactions.csv'
recipes = pd.read_csv(recipes_fp)
reviews = pd.read_csv(review_fp)

temp_df = recipes.merge(reviews, how='left', left_on='id', right_on='recipe_id')
temp_df['rating'].replace(0.0, np.nan, inplace=True)
#SAY WHY YOU DID IN WEBSITE

#average rating per recipe
```

```

avg_rating = temp_df.groupby('id').mean()[['rating']].reset_index()
avg_rating

#final dataframe
df = temp_df.merge(avg_rating, how = 'left', on = 'id')
df = df.rename(columns = {'rating_y': 'avg_rating', 'rating_x': 'rating'})

```

```

[ ]: #COLUMN DESCRIPTIONS
recipes_desc = {
    'id' : 'The recipe ID, which is unique per recipe',
    'minutes' : 'The number of minutes it takes to complete a recipe',
    'nutrition' : 'a string (that looks like a list) of various nutrition_
↳facts including calories (#), total fat (PDV), sugar (PDV), sodium (PDV),_
↳protein (PDV), saturated fat (PDV), carbohydrates (PDV)',
    'n_steps' : 'the number of steps in the recipe',
    'ingredients' : 'a string (that looks like a list) of ingredients used_
↳in the recipe',
    'n_ingredients' : 'the number of ingredients used in the recipe'
}
reviews_desc = {
    'recioe_id': 'the recipe id, matching the id column in the recipes_
↳dataframe',
    'rating' : 'the rating given by the reviewer on a 1 - 5 scale'
}

recipes_desc_df = pd.DataFrame(list(recipes_desc.items()), columns=['Column',_
↳'Description'])
reviews_desc_df = pd.DataFrame(list(reviews_desc.items()), columns=['Column',_
↳'Description'])
reviews_desc_df

```

```

[ ]:      Column      Description
0  recioe_id  the recipe id, matching the id column in the r...
1    rating  the rating given by the reviewer on a 1 - 5 scale

```

```

[ ]: recipes_html_table = recipes_desc_df.to_html(index=False)
reviews_html_table = reviews_desc_df.to_html(index=False)
with open('recipes_descriptions.html', 'w') as f:
    f.write(recipes_html_table)
with open('reviews_descriptions.html', 'w') as f:
    f.write(reviews_html_table)

```

```

[ ]: # Question: how does the length of the recipe/ ingredients affect the ratings_
↳of the recipe?

```

1.2 Step 2: Data Cleaning and Exploratory Data Analysis

```
[ ]: #split up the nutrition column to have each value in the list have its own
      ↳column and make then floats so easier to work with and pull numbers for
      ↳analysis
new_df = df.copy()

new_df['nutrition'] = new_df['nutrition'].astype(str).str.strip('[]')
new_df['nutrition'] = new_df['nutrition'].str.split(', ').apply(lambda x:
      ↳[float(i) for i in x])

def split_list(row):
    return pd.Series(row['nutrition'])

nutri = new_df[['nutrition']]
nutri = new_df.apply(split_list, axis=1)
nutri.columns = ['calories (#)', 'total fat (PDV)', 'sugar (PDV)', 'sodium
      ↳(PDV)', 'protein (PDV)', 'saturated fat (PDV)', 'carbohydrates (PDV)']
new_df = pd.concat([new_df, nutri], axis=1).drop(columns = ['nutrition',
      ↳'recipe_id', 'review'])
new_df['has_sugar'] = new_df['ingredients'].apply(lambda ingredients: 'sugar'
      ↳in ingredients)
```

```
[ ]: new_df
```

```
[ ]:
```

		name	id	minutes	\
0	1	brownies in the world best ever	333281	40	
1	1	in canada chocolate chip cookies	453467	45	
2	412	broccoli casserole	306168	40	
3	412	broccoli casserole	306168	40	
4	412	broccoli casserole	306168	40	
...		
234424		zydeco ya ya deviled eggs	308080	40	
234425		cookies by design cookies on a stick	298512	29	
234426	cookies by design	sugar shortbread cookies	298509	20	
234427	cookies by design	sugar shortbread cookies	298509	20	
234428	cookies by design	sugar shortbread cookies	298509	20	

	contributor_id	submitted	\
0	985201	2008-10-27	
1	1848091	2011-04-11	
2	50969	2008-05-30	
3	50969	2008-05-30	
4	50969	2008-05-30	
...	
234424	37779	2008-06-07	
234425	506822	2008-04-15	

234426	506822	2008-04-15
234427	506822	2008-04-15
234428	506822	2008-04-15

	tags	n_steps	\
0	['60-minutes-or-less', 'time-to-make', 'course...	10	
1	['60-minutes-or-less', 'time-to-make', 'cuisin...	12	
2	['60-minutes-or-less', 'time-to-make', 'course...	6	
3	['60-minutes-or-less', 'time-to-make', 'course...	6	
4	['60-minutes-or-less', 'time-to-make', 'course...	6	
...	
234424	['60-minutes-or-less', 'time-to-make', 'course...	7	
234425	['30-minutes-or-less', 'time-to-make', 'course...	9	
234426	['30-minutes-or-less', 'time-to-make', 'course...	5	
234427	['30-minutes-or-less', 'time-to-make', 'course...	5	
234428	['30-minutes-or-less', 'time-to-make', 'course...	5	

	steps	\
0	['heat the oven to 350f and arrange the rack i...	
1	['pre-heat oven the 350 degrees f', 'in a mixi...	
2	['preheat oven to 350 degrees', 'spray a 2 qua...	
3	['preheat oven to 350 degrees', 'spray a 2 qua...	
4	['preheat oven to 350 degrees', 'spray a 2 qua...	
...	...	
234424	['in a bowl , combine the mashed yolks and may...	
234425	['place melted butter in a large mixing bowl a...	
234426	['whip sugar and shortening in a large bowl , ...	
234427	['whip sugar and shortening in a large bowl , ...	
234428	['whip sugar and shortening in a large bowl , ...	

	description	\
0	these are the most; chocolatey, moist, rich, d...	
1	this is the recipe that we use at my school ca...	
2	since there are already 411 recipes for brocco...	
3	since there are already 411 recipes for brocco...	
4	since there are already 411 recipes for brocco...	
...	...	
234424	deviled eggs, cajun-style	
234425	i've heard of the 'cookies by design' company,...	
234426	i've heard of the 'cookies by design' company,...	
234427	i've heard of the 'cookies by design' company,...	
234428	i've heard of the 'cookies by design' company,...	

	ingredients	...	rating	\
0	['bittersweet chocolate', 'unsalted butter', '...	...	4.0	
1	['white sugar', 'brown sugar', 'salt', 'margar...	...	5.0	
2	['frozen broccoli cuts', 'cream of chicken sou...	...	5.0	

```

3      ['frozen broccoli cuts', 'cream of chicken sou...  ...  5.0
4      ['frozen broccoli cuts', 'cream of chicken sou...  ...  5.0
...
234424 ['hard-cooked eggs', 'mayonnaise', 'dijon must...  ...  5.0
234425 ['butter', 'eagle brand condensed milk', 'ligh...  ...  1.0
234426 ['granulated sugar', 'shortening', 'eggs', 'fl...  ...  1.0
234427 ['granulated sugar', 'shortening', 'eggs', 'fl...  ...  5.0
234428 ['granulated sugar', 'shortening', 'eggs', 'fl...  ...  NaN

```

	avg_rating	calories (#)	total fat (PDV)	sugar (PDV)	sodium (PDV)	\
0	4.0	138.4	10.0	50.0	3.0	
1	5.0	595.1	46.0	211.0	22.0	
2	5.0	194.8	20.0	6.0	32.0	
3	5.0	194.8	20.0	6.0	32.0	
4	5.0	194.8	20.0	6.0	32.0	
...	
234424	5.0	59.2	6.0	2.0	3.0	
234425	1.0	188.0	11.0	57.0	11.0	
234426	3.0	174.9	14.0	33.0	4.0	
234427	3.0	174.9	14.0	33.0	4.0	
234428	3.0	174.9	14.0	33.0	4.0	

	protein (PDV)	saturated fat (PDV)	carbohydrates (PDV)	has_sugar
0	3.0	19.0	6.0	True
1	13.0	51.0	26.0	True
2	22.0	36.0	3.0	False
3	22.0	36.0	3.0	False
4	22.0	36.0	3.0	False
...
234424	6.0	5.0	0.0	False
234425	7.0	21.0	9.0	True
234426	4.0	11.0	6.0	True
234427	4.0	11.0	6.0	True
234428	4.0	11.0	6.0	True

[234429 rows x 23 columns]

```

[ ]: df_html = new_df[['name', 'id', 'minutes', 'n_steps', 'n_ingredients', '
      ↪ calories (#)', 'has_sugar', 'avg_rating']].head().to_html(index=False)
with open('df_head.html', 'w') as f:
    f.write(df_html)

```

```

[ ]: # DISTRIBUTION OF NUMBER OF STEPS
num_steps = new_df.copy()
num_steps = num_steps.drop_duplicates(subset='id')
fig = px.histogram(num_steps, x='n_steps', nbins=30,
                  title='Distribution of Number of Steps in Recipe',

```

```

        labels={'n_steps': 'Number of Steps', 'count': 'Frequency'})
# Show the plot
fig.show()

```

```
[ ]: fig.write_html('univariate-n_steps.html', include_plotlyjs='cdn')
```

```
[ ]: num_steps['n_steps'].describe()
Q1 = new_df['n_steps'].quantile(0.25)
Q3 = new_df['n_steps'].quantile(0.75)

# Step 2: Calculate IQR
IQR = Q3 - Q1

# Step 3: Determine outlier boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
upper_bound

```

```
[ ]: 23.5
```

```
[ ]: #MINUTES DISTRIBUTION

minutes = new_df.copy()
minutes = minutes[minutes['minutes'] < 120]

#The distribution of minutes is skewed to the right with a very long tail.
↳Because of this, I decided to only show values where the time taken was less
↳than 120 minutes, or 2 hours,
#as I ruled out the other values that are considered outliers using the IQR
↳test. This allows us to take a better look at the distribution of the
↳minutes for most of the data
Q1 = new_df['minutes'].quantile(0.25)
Q3 = new_df['minutes'].quantile(0.75)

# Step 2: Calculate IQR
IQR = Q3 - Q1

# Step 3: Determine outlier boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

#Step 4: Filter out the outliers
filtered_minutes = new_df[(new_df['minutes'] >= lower_bound) &
↳(new_df['minutes'] <= upper_bound)]

```

```
fig = px.histogram(filtered_minutes, x='minutes', nbins=30,
                    title='Distribution of Time Taken for Recipe (minutes)',
                    labels={'log_calories': 'Minutes', 'count': 'Frequency'})

# Show the plot
fig.show()
```

```
[ ]: fig.write_html('univariate-minutes.html', include_plotlyjs='cdn')
```

```
[ ]: #NUMBER OF STEPS VS NUMBER OF INGREDIENTS
step_ingredients = new_df.groupby('id').mean()#[['n_steps', 'n_ingredients']]
fig = px.scatter(step_ingredients, x='avg_rating', y='n_steps',
                 title='Average Rating vs. Number of Steps',
                 labels={'avg_rating': 'Average Rating', 'n_steps': 'Number of Steps'})

# Show the plot
fig.show()
```

```
[ ]: fig.write_html('n_steps-rating.html', include_plotlyjs='cdn')
```

```
[ ]: calories_ingredients = new_df.groupby('id').mean()#[['n_steps', 'n_ingredients']]
fig = px.scatter(step_ingredients, x='n_ingredients', y='n_steps',
                 title='Number of Calories vs. Number of Steps',
                 labels={'avg_rating': 'Average Rating', 'n_steps': 'Number of Steps'})

# Show the plot
fig.show()
```

```
[ ]: Q1 = new_df['calories (#)'].quantile(0.25)
Q3 = new_df['calories (#)'].quantile(0.75)

# Step 2: Calculate IQR
IQR = Q3 - Q1

# Step 3: Determine outlier boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

#Step 4: Filter out the outliers
filtered_calories = new_df[(new_df['calories (#)'] >= lower_bound) &
                             (new_df['calories (#)'] <= upper_bound)]

pivot_table = filtered_calories[['calories (#)', 'n_ingredients']].
    pivot_table(index = 'n_ingredients', aggfunc = ['mean', 'median', 'std'])
pivot_table
```

```
[ ]:
```

	mean	median	std
	calories (#)	calories (#)	calories (#)
n_ingredients			
1	157.229630	144.20	113.179463
2	212.604051	144.70	205.470241
3	215.911121	163.60	191.614204
4	237.279823	183.60	193.219294
5	262.716944	219.00	191.958962
6	285.390661	239.90	207.755204
7	302.350313	252.20	205.388476
8	318.771192	275.20	204.539190
9	336.953357	294.90	207.894473
10	341.251078	304.10	201.831227
11	363.584400	324.80	207.621016
12	364.438348	326.40	202.957711
13	384.200156	359.50	205.592255
14	402.556815	374.30	211.273798
15	422.722760	392.50	215.983254
16	433.215455	416.70	210.721340
17	458.555769	419.70	216.732349
18	469.239942	440.20	215.431545
19	469.590974	439.10	224.048137
20	498.537207	452.70	237.817307
21	452.178998	416.80	218.114314
22	572.306952	620.50	242.775594
23	476.350549	471.10	211.118148
24	505.260976	454.50	214.630693
25	540.551163	606.40	271.928627
26	544.987273	572.70	178.086653
27	583.072222	589.85	214.616784
28	559.124324	491.70	257.784868
29	442.145000	336.20	169.228984
30	580.293103	594.30	165.961585
31	348.242857	219.60	172.555632
32	363.100000	363.10	NaN
33	338.200000	338.20	NaN

```
[ ]: html_pivot = pivot_table.to_html()
with open('pivot_table.html', 'w') as f:
    f.write(html_pivot)
```

1.3 Step 3: Assessment of Missingness

```
[ ]: # TODO
nan_counts = df.isna().sum()
nan_counts
```



```
[ ]: name          1
      id           0
      minutes      0
      contributor_id 0
      submitted    0
      tags         0
      nutrition    0
      n_steps      0
      steps        0
      description  114
      ingredients  0
      n_ingredients 0
      user_id      1
      recipe_id    1
      date         1
      rating       15036
      review       58
      avg_rating   2777
      dtype: int64
```

```
[ ]: #PERMUTATION TEST TO ASSESS MISSINGNESS OF REVIEW
      #RATING VS. MINUTES

      #null hypothesis: missingness of rating does not depend on the number of minutes
      #alternative hypothesis: the missingness of the rating depends on the number of
      minutes

      def stat(avg_rating, minutes):
          is_missing = avg_rating.isna()
          mean_missing = minutes[is_missing].mean()
          mean_notmissing = minutes[~is_missing].mean()

          return np.abs(mean_notmissing - mean_missing)

      #make a copy of the og dataframe, drop the duplicates of the recipe because the
      description and number of steps doesnt change within a recipe
      shuffled = new_df.copy().drop_duplicates(subset=['id'])

      observed = stat(shuffled['avg_rating'], shuffled['minutes'])

      stats = []

      for i in range(1000):
          shuffled['shuffled rating'] = np.random.permutation(shuffled['avg_rating'])
          curr = stat(shuffled['shuffled rating'], shuffled['minutes'])
          stats = np.append(stats, curr)
```

```

p_value = np.mean(np.array(stats) >= observed)
p_value

# we get a p-value of 0.035 and there fore reject the null hypothesis at the 0.
↳01 significance level

```

```
[ ]: 0.039
```

```

[ ]: fig = px.histogram(
    pd.DataFrame(stats), x=0, nbins=10, histnorm='probability',
    title='Empirical Distribution of Mean Differences in Minutes (Missing vs.↳
↳Not Missing Ratings) ')
fig.add_vline(x=observed, line_color='red')
fig.update_xaxes(title_text='Mean Differences in Minutes')

```

```
[ ]: fig.write_html('missing_test.html', include_plotlyjs='cdn')
```

```

[ ]: #null hypothesis: missingness of rating does not depend on the number of↳
↳calories in the recipe
#alternative hypothesis: the missingness of the rating depends on the number of↳
↳calories in the recipe

def stat(avg_rating, calories):
    is_missing = avg_rating.isna()
    mean_missing = calories[is_missing].mean()
    mean_notmissing = calories[~is_missing].mean()

    return np.abs(mean_notmissing - mean_missing)

#make a copy of the og dataframe, drop the duplicates of the recipe because the↳
↳description and number of steps doesnt change within a recipe
shuffled = new_df.copy().drop_duplicates(subset=['id'])

observed = stat(shuffled['avg_rating'], shuffled['calories (#)'])

stats = []

for i in range(1000):
    shuffled['shuffled rating'] = np.random.permutation(shuffled['avg_rating'])
    curr = stat(shuffled['shuffled rating'], shuffled['calories (#)'])
    stats = np.append(stats, curr)

p_value = np.mean(np.array(stats) >= observed)
p_value

#we get a p-value of 0.0, and therefore reject the null hypothesis at a 0.01↳
↳significant level

```

```
[ ]: 0.0
```

1.4 Step 4: Hypothesis Testing

```
[ ]: # TODO
#null hyp: there is no relationship between time it takes for a recipe and the
    ↳ average rating of a recipe
#alt hype: recipes that take over 37 minutes have a lower average rating than
    ↳ ones that take 37 or less minutes

# #FIX UP DATAFRAME
# #remove outliers
Q1 = new_df['minutes'].quantile(0.25)
Q3 = new_df['minutes'].quantile(0.75)

# Step 2: Calculate IQR
IQR = Q3 - Q1

# Step 3: Determine outlier boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

#Step 4: Filter out the outliers
hyp_test_df = new_df[(new_df['minutes'] >= lower_bound) & (new_df['minutes'] <=
    ↳ upper_bound)]

# #keep only necessary columns and remove duplicates because we only need one
    ↳ for each recipe, as the minutes and average rating wont change
# hyp_test_df = new_df[['id', 'minutes', 'avg_rating']].drop_duplicates()
hyp_test_df = new_df.copy()

[ ]: hyp_test_df['under 37'] = hyp_test_df['minutes'] <=37
hyp_test_df
```

```
[ ]:
```

		name	id	minutes	\
0	1 brownies in the world	best ever	333281	40	
1	1 in canada chocolate chip cookies		453467	45	
2	412 broccoli casserole		306168	40	
3	412 broccoli casserole		306168	40	
4	412 broccoli casserole		306168	40	
...		
234424	zydeco ya ya deviled eggs		308080	40	
234425	cookies by design cookies on a stick		298512	29	
234426	cookies by design sugar shortbread cookies		298509	20	
234427	cookies by design sugar shortbread cookies		298509	20	
234428	cookies by design sugar shortbread cookies		298509	20	

	contributor_id	submitted \
0	985201	2008-10-27
1	1848091	2011-04-11
2	50969	2008-05-30
3	50969	2008-05-30
4	50969	2008-05-30
...
234424	37779	2008-06-07
234425	506822	2008-04-15
234426	506822	2008-04-15
234427	506822	2008-04-15
234428	506822	2008-04-15

	tags	n_steps \
0	['60-minutes-or-less', 'time-to-make', 'course...	10
1	['60-minutes-or-less', 'time-to-make', 'cuisin...	12
2	['60-minutes-or-less', 'time-to-make', 'course...	6
3	['60-minutes-or-less', 'time-to-make', 'course...	6
4	['60-minutes-or-less', 'time-to-make', 'course...	6
...
234424	['60-minutes-or-less', 'time-to-make', 'course...	7
234425	['30-minutes-or-less', 'time-to-make', 'course...	9
234426	['30-minutes-or-less', 'time-to-make', 'course...	5
234427	['30-minutes-or-less', 'time-to-make', 'course...	5
234428	['30-minutes-or-less', 'time-to-make', 'course...	5

	steps \
0	['heat the oven to 350f and arrange the rack i...
1	['pre-heat oven the 350 degrees f', 'in a mixi...
2	['preheat oven to 350 degrees', 'spray a 2 qua...
3	['preheat oven to 350 degrees', 'spray a 2 qua...
4	['preheat oven to 350 degrees', 'spray a 2 qua...
...	...
234424	['in a bowl , combine the mashed yolks and may...
234425	['place melted butter in a large mixing bowl a...
234426	['whip sugar and shortening in a large bowl , ...
234427	['whip sugar and shortening in a large bowl , ...
234428	['whip sugar and shortening in a large bowl , ...

	description \
0	these are the most; chocolatey, moist, rich, d...
1	this is the recipe that we use at my school ca...
2	since there are already 411 recipes for brocco...
3	since there are already 411 recipes for brocco...
4	since there are already 411 recipes for brocco...
...	...
234424	deviled eggs, cajun-style

234425 i've heard of the 'cookies by design' company,...
 234426 i've heard of the 'cookies by design' company,...
 234427 i've heard of the 'cookies by design' company,...
 234428 i've heard of the 'cookies by design' company,...

	ingredients	avg_rating	\
0	['bittersweet chocolate', 'unsalted butter', '... ..	4.0	
1	['white sugar', 'brown sugar', 'salt', 'margar... ..	5.0	
2	['frozen broccoli cuts', 'cream of chicken sou... ..	5.0	
3	['frozen broccoli cuts', 'cream of chicken sou... ..	5.0	
4	['frozen broccoli cuts', 'cream of chicken sou... ..	5.0	
...	
234424	['hard-cooked eggs', 'mayonnaise', 'dijon must... ..	5.0	
234425	['butter', 'eagle brand condensed milk', 'ligh... ..	1.0	
234426	['granulated sugar', 'shortening', 'eggs', 'fl... ..	3.0	
234427	['granulated sugar', 'shortening', 'eggs', 'fl... ..	3.0	
234428	['granulated sugar', 'shortening', 'eggs', 'fl... ..	3.0	

	calories (#)	total fat (PDV)	sugar (PDV)	sodium (PDV)	\
0	138.4	10.0	50.0	3.0	
1	595.1	46.0	211.0	22.0	
2	194.8	20.0	6.0	32.0	
3	194.8	20.0	6.0	32.0	
4	194.8	20.0	6.0	32.0	
...	
234424	59.2	6.0	2.0	3.0	
234425	188.0	11.0	57.0	11.0	
234426	174.9	14.0	33.0	4.0	
234427	174.9	14.0	33.0	4.0	
234428	174.9	14.0	33.0	4.0	

	protein (PDV)	saturated fat (PDV)	carbohydrates (PDV)	has_sugar	\
0	3.0	19.0	6.0	True	
1	13.0	51.0	26.0	True	
2	22.0	36.0	3.0	False	
3	22.0	36.0	3.0	False	
4	22.0	36.0	3.0	False	
...	
234424	6.0	5.0	0.0	False	
234425	7.0	21.0	9.0	True	
234426	4.0	11.0	6.0	True	
234427	4.0	11.0	6.0	True	
234428	4.0	11.0	6.0	True	

under 37

0	False
1	False

```

2         False
3         False
4         False
...
234424    False
234425     True
234426     True
234427     True
234428     True

```

```
[234429 rows x 24 columns]
```

```
[ ]: observed = hyp_test_df.groupby('under 37')['avg_rating'].mean()
observed_diff = observed[True] - observed[False]

reps = 1000
diffs = []

for i in range(reps):
    with_shuffled = hyp_test_df.assign(shuffled_rating=np.random.
    ↪ permutation(hyp_test_df['avg_rating']))
    group_means = (with_shuffled.groupby('under 37').mean().loc[:
    ↪ , 'shuffled_rating'])
    diff = group_means.loc[True] - group_means.loc[False]
    diffs.append(diff)
```

```
[ ]: fig = px.histogram(
    pd.DataFrame(diffs), x=0, nbins=10, histnorm='probability',
    title='Empirical Distribution of the Mean Differences <br> in Average_
    ↪ Rating (Recipes Under 37 Minutes - Over 37 Minutes)')
fig.add_vline(x=observed_diff, line_color='red')
fig.update_layout(xaxis_range=[-0.03, 0.04], margin=dict(t=60))
fig.update_xaxes(title_text='Mean Differences in Rating')
```

```
[ ]: fig.write_html('hyp_test.html', include_plotlyjs='cdn')
```

```
[ ]: p_value = np.mean(np.array(diffs) >= observed_diff)
p_value
```

```
[ ]: 0.0
```

1.5 Step 5: Framing a Prediction Problem

```
[ ]: # TODO

#The prediction problem I will focus on will be predicting the rating of a
    ↪ recipe. For this, I will use a Random Forest Classifier,
```

```
# as the rating can be considered as categorical when I round the average
↳ rating column to be an integer
```

1.6 Step 6: Baseline Model

```
[ ]: # TODO
new_df.columns
#remove columns that have na in the average col. means that there was no rating
↳ on that recipe
new_rating = new_df[~new_df['avg_rating'].isna()]

new_rating = new_rating[['n_steps', 'n_ingredients', 'avg_rating']]
new_rating['rounded_rating'] = new_rating['avg_rating'].round().astype(int)
new_rating = new_rating.drop(columns=['avg_rating'])
new_rating
```

```
[ ]:
```

	n_steps	n_ingredients	rounded_rating
0	10	9	4
1	12	11	5
2	6	9	5
3	6	9	5
4	6	9	5
...
234424	7	8	5
234425	9	10	1
234426	5	7	3
234427	5	7	3
234428	5	7	3

```
[231652 rows x 3 columns]
```

```
[ ]: X_train, X_test, y_train, y_test = (
    train_test_split(new_rating.drop(columns=['rounded_rating']),
        ↳ new_rating['rounded_rating'], random_state=1)
)
```

```
[ ]: preprocessor = ColumnTransformer(
    transformers=[
        ('std-scalar', StandardScaler(), ['n_steps', 'n_ingredients'])
    ]
)
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', RandomForestClassifier(random_state=42))
])
pipeline.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('std-scalar',
                                                       StandardScaler(),
                                                       ['n_steps',
                                                        'n_ingredients'])])),
                    ('classifier', RandomForestClassifier(random_state=42))])
```

```
[ ]: y_pred = pipeline.predict(X_test)
```

```
[ ]: f1_score(y_test, y_pred, average='weighted')
```

```
[ ]: 0.6386647045521429
```

1.7 Step 7: Final Model

```
[ ]: fin_mod = new_df.copy()
fin_mod = fin_mod[~fin_mod['avg_rating'].isna()]
fin_mod = fin_mod[['avg_rating', 'minutes', 'n_steps', 'n_ingredients', '
↳ calories (#)', 'has_sugar']]
fin_mod['rounded_rating'] = fin_mod['avg_rating'].round().astype(int)
fin_mod = fin_mod.drop(columns=['avg_rating'])
fin_mod
```

```
[ ]:
```

	minutes	n_steps	n_ingredients	calories (#)	has_sugar	\
0	40	10	9	138.4	True	
1	45	12	11	595.1	True	
2	40	6	9	194.8	False	
3	40	6	9	194.8	False	
4	40	6	9	194.8	False	
...	
234424	40	7	8	59.2	False	
234425	29	9	10	188.0	True	
234426	20	5	7	174.9	True	
234427	20	5	7	174.9	True	
234428	20	5	7	174.9	True	

```
rounded_rating
```

0	4
1	5
2	5
3	5
4	5
...	...
234424	5
234425	1
234426	3
234427	3

234428

3

[231652 rows x 6 columns]

```
[ ]: # TODO
#transform calories (#) with standard scalar
#one hot encode ingredients
#one hot encode tags
#transform minutes with standard scalar

X = fin_mod.drop(columns = ['rounded_rating'])
y = fin_mod['rounded_rating']
X_train, X_test, y_train, y_test = (
    train_test_split(X, y, random_state=1)
)

preproc = ColumnTransformer(
    transformers = [
        ('std-scale', StandardScaler(), ['calories (#)', 'minutes']),
        ('one-hot', OneHotEncoder(drop='first', handle_unknown='ignore'),
        ['has_sugar']),
    ],
    remainder= 'passthrough'
)
pl = Pipeline([
    ('preprocessor', preproc),
    ('forest', RandomForestClassifier(max_depth = 3, n_estimators = 50))
])

#FIND THE CORRECT HYPERPARAMETERS
param_grid = {
    'forest_max_depth': [3, 5, 7, 10, None],
    'forest_n_estimators': [50, 100, 200]
}

grid_search = GridSearchCV(pl, param_grid, cv=5, scoring='accuracy', n_jobs=1)
grid_search.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('preprocessor',
ColumnTransformer(remainder='passthrough',
                                transformers=[('std-
scale',
StandardScaler(),
['calories '
                                ' (#)',
```

```

'minutes'])),
                                                                    ('one-
hot',
OneHotEncoder(drop='first',
               handle_unknown='ignore'),
['has_sugar']]))),
                                                                    ('forest',
RandomForestClassifier(max_depth=3,
n_estimators=50))),
               n_jobs=1,
               param_grid={'forest__max_depth': [3, 5, 7, 10, None],
                           'forest__n_estimators': [50, 100, 200]},
               scoring='accuracy')

```

```

[ ]: best_params = grid_search.best_params_
best_score = grid_search.best_score_

best_params

```

```

[ ]: {'forest__max_depth': None, 'forest__n_estimators': 200}

```

```

[ ]: y_pred = grid_search.predict(X_test)

```

```

[ ]: f1_score(y_test, y_pred, average='weighted')

```

```

[ ]: 0.9133475035483493

```

1.8 Step 8: Fairness Analysis

```

[ ]: # TODO
#split between recipes over 35 minutes and under 35 minutes
#run permutation test to check difference in precision between two groups
#null hyp: the model is fair. The precision for shorter recipes is roughly the
    ↳ same as longer recipes
#alt hyp: the model is unfair. The precision for shorter recipes is lower than
    ↳ longer recipes

```

```

[ ]: new_df['minutes'].median()

```

```

[ ]: 35.0

```

```

[ ]: shorter = fin_mod[fin_mod['minutes'] <= 35]
shorter_y = shorter['rounded_rating']
shorter_x = shorter.drop(columns = ['rounded_rating'])
shorter_pred = grid_search.predict(shorter_x)
shorter_prediction = precision_score(shorter_y, shorter_pred,
    ↳ average='weighted')

```

```
shorter_prediction
```

```
[ ]: 0.9792886999818651
```

```
[ ]: longer = fin_mod[fin_mod['minutes'] > 35]
longer_y = longer['rounded_rating']
longer_x = longer.drop(columns = ['rounded_rating'])
longer_pred = grid_search.predict(longer_x)
longer_prediction = precision_score(longer_y, longer_pred, average='weighted')
longer_prediction
```

```
[ ]: 0.9788856657967512
```

```
[ ]: observed = shorter_prediction - longer_prediction
observed
```

```
[ ]: 0.000403034185113893
```

```
[ ]: #permutation test
diffs = []
precision_df = fin_mod.copy()
```

```
[ ]: for i in range(100):
    shuffled = np.random.permutation(precision_df['rounded_rating'])
    precision_df = precision_df.assign(shuffled_rating = shuffled)

    shorter = precision_df[precision_df['minutes'] <= 35]
    shorter_y = shorter['shuffled_rating']
    shorter_x = shorter.drop(columns = ['shuffled_rating', 'rounded_rating'])
    shorter_pred = grid_search.predict(shorter_x)
    shorter_prediction = precision_score(shorter_y, shorter_pred,
    ↪average='weighted')

    longer = precision_df[precision_df['minutes'] > 35]
    longer_y = longer['shuffled_rating']
    longer_x = longer.drop(columns = ['shuffled_rating', 'rounded_rating'])
    longer_pred = grid_search.predict(longer_x)
    longer_prediction = precision_score(longer_y, longer_pred,
    ↪average='weighted')

    diffs.append(shorter_prediction - longer_prediction)
```

```
[ ]: diffs
```

```
[ ]: [-0.0016998339853376843,
-0.0034367070983210013,
-0.0005883149926669828,
```

-0.0022284249237168874,
-0.0027897479966420002,
-0.001078699308566633,
0.0004479573064095632,
-0.0016253886811437024,
-0.003105552859107963,
-0.005947199475976106,
-0.0002718036995270623,
0.0022741938552789387,
-0.0028210665538923596,
0.0004650937305172853,
0.0024180220464607993,
0.0037865806552470627,
0.005655988940146184,
-0.00042169587053553226,
-0.000714131958881814,
0.0035712826804881193,
0.0009404385836205842,
-0.0010053325669157065,
-3.408168565377512e-05,
0.0006741225340823886,
-0.003669073336166506,
0.0013772367314695,
-0.0017880815787645332,
-0.004516893834347835,
-0.0033615726863786,
0.0025014192168028027,
0.0004240023547243954,
0.002648199668868312,
0.0007866152003118687,
0.004613735020790877,
-0.0013676838275579195,
0.0028710705372679834,
0.0008992945629394677,
0.0014695961740602836,
0.0009339743577975179,
-9.718639657463335e-05,
-0.001766939802653078,
-0.001333015934697368,
0.0027650655015563075,
-0.004749897251573376,
-0.0026602914799384036,
0.002859615652082814,
-0.002145991081213494,
-0.0014358815111293888,
0.0008729149426810467,
0.003563366465701878,

0.0024910337518007086,
-0.0024335211613922825,
0.0008458757452383114,
0.004326250539573473,
0.0014667774167907988,
0.002206024450906341,
0.0010228784865179419,
-0.00021576026878245003,
-0.0029616605824426268,
-0.001559995383600854,
-0.0023563412786240523,
-0.0015365031994247769,
0.0009453699304010632,
-0.0034655224087873915,
0.004441390572462622,
-0.00369896855483709,
-0.001864270810740476,
0.0007865923628006533,
-0.0008794365861319875,
0.004138231193082542,
0.0018301860594152064,
0.0017263235794487963,
-0.004624106255835869,
-0.00046171471217171245,
-0.0016752134984703293,
0.001845591446487438,
0.002409355676392555,
0.001719731601813268,
0.001484121061421817,
-0.0014363030173651925,
0.0005326241544622023,
0.0018881591354553695,
-0.0013406083118409073,
-0.001491404182760947,
-0.0008065737031349718,
-0.0014488471832587724,
-0.00028051794522354623,
-0.0004510262879324767,
0.003579459252160966,
0.0018446227783069924,
-0.0032880719350415477,
0.0013483530758955364,
0.0047582224497441095,
-0.001695151678702711,
-0.001604868132056514,
0.0034620593461596183,
-0.0012418937713889466,

```
-0.001738371263766858,  
-0.00463602531163454,  
-2.206453360742433e-05]
```

```
[ ]: p_value = np.mean(np.abs(diffs) >= np.abs(observed))  
p_value
```

```
[ ]: 0.94
```

```
[ ]: fig = px.histogram(  
    pd.DataFrame(diffs), x=0, nbins=10, histnorm='probability',  
    title='Empirical Distribution of the Precision Differences <br> in Shorter_␣  
    ↪vs Longer Recipes (shorter - longer)')  
fig.add_vline(x=observed, line_color='red')  
fig.update_layout(xaxis_range=[-0.01, 0.01], margin=dict(t=60))  
fig.update_xaxes(title_text= 'Difference in Precision')
```

```
[ ]: fig.write_html('fairness.html', include_plotlyjs='cdn')
```

```
[ ]:
```